

A Review of Sensitivity Methods for Differential Equations

Facundo Sapienza^{*1}, Jordi Bolibar², Frank Schäfer³, Patrick Heimbach⁶, Giles Hooker⁴, Fernando Pérez¹, and Per-Olof Persson⁵

¹*Department of Statistics, University of California, Berkeley (USA)*

²*TU Delft, Department of Geosciences and Civil Engineering, Delft (Netherlands)*

³*CSAIL, Massachusetts Institute of Technology, Cambridge (USA)*

⁴*Department of Statistics and Data Science, University of Pennsylvania (USA)*

⁵*Department of Mathematics, University of California, Berkeley (USA)*

⁶*Department of Earth and Planetary Sciences, Jackson School of Geosciences, University of Texas, Austin (USA)*

November 14, 2023

^{*}Corresponding author: fsapienza@berkeley.edu

Abstract

Differentiable programming has become a central component of modern machine learning techniques. In the context of models described by differential equations, the calculation of sensitivities and gradients requires careful algebraic and numeric manipulations of the underlying dynamical system. We aim to summarize some of the most used techniques that exist to compute gradients for numerical solutions of differential equations. We cover this problem by first introducing motivations in current areas of research, such as geophysics; the mathematical foundations of the different approaches; and finally the computational consideration and solutions that exist in modern scientific software.

To the community, by the community. *This manuscript was conceived with the goal of shortening the gap between developers and practitioners of differentiable programming applied to modern scientific machine learning. With the advent of new tools and new software, it is important to create pedagogical content that allows the broader community to understand and integrate these methods into their workflows. We hope this encourages new people to be an active part of the ecosystem, by using and developing open-source tools. This work was done under the premise **open-science from scratch**, meaning all the contents of this work, both code and text, have been in the open from the beginning and that any interested person can contribute to the project. You can contribute directly to the GitHub repository github.com/ODINN-SciML/DiffEqSensitivity-Review.*

Contents

1	Introduction	4
2	Scientific motivation	6
2.1	On the importance of differentiable programming	6
2.2	Domain-specific applications	7
2.2.1	Computational Fluid Dynamics	7
2.2.2	Electromagnetism	7
2.2.3	Geosciences	7
2.2.4	Biology and Ecology	8
2.2.5	Quantum Physics	8
3	Methods	8
3.1	Preliminaries	9
3.2	Finite differences	11
3.3	Complex step differentiation	12
3.4	Automatic differentiation	12
3.4.1	Forward mode	13
3.4.2	Backward mode	14
3.4.3	AD connection with JVPs and VJPs	15
3.5	Symbolic differentiation	16
3.6	Sensitivity equations	17
3.7	Adjoint methods	17
3.7.1	Discrete adjoint method	18
3.7.2	Continuous adjoint method	21
4	Computational implementation	22
4.1	Direct methods	22
4.1.1	Finite differences	23
4.1.2	Forward-mode AD	23
4.1.3	Reverse-mode AD	25
4.2	Solver-based methods	25
4.2.1	Sensitivity equation	25
4.2.2	Solving the adjoint	25
4.2.3	Further considerations	27
	Appendices	28
A	Lagrangian derivation of adjoints	28
B	Supplementaty code	28
	References	29

Plain language summary

Scientific models are used to predict and understand a vast array of different dynamics, ranging from physical processes to ecological, biological, and social interactions or chemical reactions, among many others. The combination of mechanistic models with data-driven models is becoming increasingly common in many scientific domains. In order to achieve this, these models need to leverage both domain knowledge and data, in order to have an accurate representation of the underlying dynamics. Being able to determine which model parameters are most influential and further compute derivatives of such a model is key to correctly assimilating and learning from data, but a myriad of sensitivity methods exist to do so. We provide an overview of the different sensitivity methods that exist, providing (i) guidelines on the best use cases for different scientific domain problems, (ii) detailed mathematical analyses of their characteristics, and (iii) computational implementations on how to solve them efficiently.

1 Introduction

Evaluating how the value of a function changes with respect to its arguments and parameters plays a central role in optimization, sensitivity analysis, Bayesian inference, inverse methods, and uncertainty quantification, among many. Modern machine learning applications require the use of gradients to explore and exploit more efficiently the space of parameters (e.g., weight of a neural network). When optimizing a loss function, gradient-based methods (for example, gradient descent and its many variants⁵⁴) are more efficient at finding a minimum and converge faster to them than gradient-free methods. When numerically computing the posterior of a probabilistic model, gradient-based sampling strategies converge faster to the posterior distribution than gradient-free methods. Hessians further help to improve the convergence rates of these algorithms and enable uncertainty quantification around parameter values. *A gradient serves as a compass in modern data science: it tells us in which direction in the open wide ocean of parameters we should move towards in order to increase our chances of success.*

Dynamical systems governed by differential equations are not an exception to the rule. Differential equations play a central role in describing the behaviour of systems in natural and social sciences. Some authors have recently suggested differentiable programming as the bridge between modern machine learning methods and traditional scientific models.^{17,52,59} Being able to compute gradients or sensitivities of dynamical systems opens the door to more complex models. This is very appealing in geophysical models, where there is a broad literature on physical models and a long tradition in numerical methods. The first goal of this work is to introduce some of the applications of this emerging technology and to motivate its incorporation for the modelling of complex systems in the natural and social sciences.

Question 1. *What are the scientific applications of differentiable programming for complex dynamical systems?*

Sensitivity analysis corresponds to any method aiming to calculate how much the output of a function or program changes when we vary one of the model parameters. This task is

performed in different ways by different communities when working with dynamical systems. In statistics, the sensitivity equations enable the computation of gradients of the likelihood of the model with respect to the parameters of the dynamical system, which can be later used for inference.⁵¹ In numerical analysis, sensitivities quantify how the solution of a differential equation fluctuates with respect to certain parameters. This is particularly useful in optimal control theory,²⁰ where the goal is to find the optimal value of some control (e.g. the shape of a wing) that minimizes a given loss function. In recent years, there has been an increasing interest in designing machine learning workflows that include constraints in the form of differential equations. Examples of this include methods that numerically solve differential equations,⁵⁰ as well as methods that augment and learn parts of the differential equation.^{13,49} Furthermore, numerical solvers are used as forward models in the case of neural ordinary differential equations.¹¹

However, when working with differential equations, the computation of gradients is not an easy task, both regarding the mathematical framework and software implementation involved. Except for a small set of particular cases, most differential equations require numerical methods to calculate their solution and cannot be differentiated analytically. This means that solutions cannot be directly differentiated and require special treatment if, besides the numerical solution, we also want to compute first or second-order derivatives. Furthermore, numerical solutions introduce approximation errors. These errors can be propagated and amplified during the computation of the gradient. Alternatively, there is a broad literature on numerical methods for solving differential equations. Although each method provides different guarantees and advantages depending on the use case, this means that the tools developed to compute gradients when using a solver need to be universal enough in order to be applied to all or at least to a large set of them. The second goal of this article is to review different methods that exist to achieve this goal.

Question 2. *How can I compute the gradient of a function that depends on the numerical solution of a differential equation?*

The broader set of tools known as Automatic Differentiation (AD) aims at computing derivatives by sequentially applying the chain rule to the sequence of unit operations that constitute a computer program. The premise is simple: every computer program, including a numerical solver, is ultimately an algorithm described by a chain of simple algebraic operations (addition, multiplication) that are easy to differentiate and their combination is easy to differentiate by using the chain rule. Although many modern differentiation tools use AD to some extent, there is also a family of methods that compute the gradient by relying on an auxiliary set of differential equations. We are going to refer to this family of methods as *continuous*, and we will dedicate them a special treatment in future sections to distinguish them from the discrete algorithms that resemble more to pure AD.

The differences between methods arise both from their mathematical formulation and their computational implementation. The first provides different guarantees on the method returning the actual gradient or a good approximation of it. The second involves how theory is translated to software, and what are the data structures and algorithms used to implement it. Different methods have different computational complexities depending on the number of parameters and differential equations, and these complexities are also balanced between total execution time and required memory. The third goal of this work is then to illustrate

the different strengths and weaknesses of these methods, and how to use them in modern scientific software.

Question 3. *What are the advantages and disadvantages of different differentiation methods and how can I incorporate them in my research?*

Despite the fact that these methods can be (in principle) implemented in different programming languages, here we decided to use the Julia programming language for the different examples. Julia is a recent but mature programming language that has already a large tradition in implementing packages aiming to advance differentiable programming.^{7,8}

Without aiming at making an extensive and specialized review of the field, we believe this study will be useful to other researchers working on problems that combine optimization and sensitivity analysis with differential equations. Differentiable programming is opening new ways of doing research across sciences, and we need close collaboration between domain scientists, methodological scientists, and computer scientists in order to develop successful, scalable, practical, and efficient frameworks for real-worlds applications. As we make progress in the use of these tools, new methodological questions start to emerge. How do these methods compare? How can their been improved? We also hope this paper serves as a gateway to new questions regarding advances in these methods.

2 Scientific motivation

2.1 On the importance of differentiable programming

Scientific models from many domains have often been based on mechanistic models, represented as differential equations, involving the use of numerical methods to solve them. Among many, this lead to fundamental advances in the physical sciences during the last century, with the combination of complex mathematical theories and a reduced amount of observations to validate them. Nonetheless, in the 21st century, with the unstoppable wave of data flooding all scientific domains, progress with such traditional methods has become more complex.

Alternatively, the field of statistics experienced a boom following the massive growth of data, signaling the era of data science and machine learning. With the advent of machine learning methods, it is possible to learn and capture extremely complex nonlinear patterns and information hidden in huge datasets. Machine learning models can be seen as the opposite of mechanistic models: they are flexible, data-driven and they do not necessarily respect domain-specific constraints.

At first sight, these two modelling philosophies can be seen as antagonistic, and this is more or less the way they have evolved in the last decades.⁶⁷ On the one hand, domain scientists have often struggled to adopt machine learning methods, judging them as opaque black boxes, unreliable, and not respecting domain-established knowledge. On the other hand, the field of machine learning has mainly been developed around data-driven applications, without including any *a priori* physical knowledge. However, there has been an increasing interest in making mechanistic models more flexible, as well as introducing domain-specific or physical constraints and interpretability in machine learning models.

A key way to achieve this is through differentiable programming, i.e. being able to compute derivatives of any computer program describing a scientific model. During the last decades, the backpropagation algorithm has enabled the fast-growing of deep learning by efficiently computing gradients of large and complex neural networks with many parameters.²⁴ Nowadays, the differentiation of hybrid models comprising data-driven models (e.g. neural networks, gaussian processes) with differential equations poses complex technical problems, which are only starting to be explored in recent years.⁴¹ Being able to accurately estimate model parameters, ranging from a few ones in classic inversion problems to millions of them in large neural networks, opens many new possibilities. Differentiable programming has the potential to revolutionize the way we approach and design scientific models and even the way we discover governing laws from data.

2.2 Domain-specific applications

Differential equations can be used to describe a large variety of dynamical systems, while data-driven regression models (e.g. neural networks, Gaussian processes, basis expansions) have been demonstrated to act as universal approximators, virtually learning any possible function if enough data is available.²² This combined flexibility can be exploited by many different domain-specific problems to tailor modelling needs to both dynamics and data characteristics.

2.2.1 Computational Fluid Dynamics

2.2.2 Electromagnetism

2.2.3 Geosciences

In geosciences, partial differential equations (PDEs) are often used to simulate fluid dynamics, describing geophysical properties of many Earth systems, such as the atmosphere, oceans, or glaciers. In such models, calibrating model parameters is extremely challenging, due to datasets being sparse in both space and time and noisy. Moreover, many existing mechanistic models can only partially describe observations, with many detailed physical processes being ignored or badly parametrized. The use of differentiable programming, combining PDEs and data-driven models (i.e. Universal Differential Equations) can add flexibility to mechanistic models in order to incorporate new governing laws from data.⁴⁹

Glaciers act as slow fluids, flowing down-slope through the effects of gravity, and the understanding of their rheological properties (e.g. ice viscosity affecting internal deformation or sliding at the glacier-bedrock interface) is key to assessing their contribution to water resources and sea-level rise.¹² These rheological processes and their dependency on key large-scale environmental variables, such as the local climate or topography, are still not well understood. The use of differentiable programming, combined with Universal Differential Equations, holds great potential to learn new empirical laws of these physical processes from large-scale remote sensing datasets. A recent study showed how Julia’s differentiable programming capabilities can be used to optimize the parameters of a neural network, learning a function of the nonlinear ice diffusivity in a glacier ice flow PDE, to match observations.⁹




Figure 1: Schematic representation of the different methods available for differentiation involving differential equation solutions. These can be classified depending if they find the gradient by solving a new system of differential equations (*continuous*) or if instead they manipulate unit algebraic operations (*discrete*). Furthermore, depending on if these methods run in the same direction as the numerical solver, we are going to be talking about *backward* and *forward* methods.

2.2.4 Biology and Ecology

2.2.5 Quantum Physics

3 Methods

Depending on the number of parameters and the complexity of the differential equation we are trying to solve, there are different methods to compute gradients with different numerical and computational advantages. These methods can be roughly classified as:

- *Discrete* vs *continuous* methods
- *Forward* vs *backward* methods

The first difference regards the fact that the method for computing the gradient can be either based on the manipulation of atomic operations that are easy to differentiate using the chain rule several times (discrete), in opposition to the approach of approximating the gradient as the numerical solution of a new set of differential equations (continuous). Another way of conceptualizing this difference is by comparing them with the discretize-differentiate and differentiate-discretize approaches.^{10,48,60,68} We can either discretize the original system of ODEs in order to numerically solve it and then define the set of adjoint equations on top of the numerical scheme; or instead define the adjoint equation directly using the differential equation and then discretize both in order to solve.²⁰

The second distinction is related to the fact that some methods compute gradients by resolving a new sequential problem that may move in the same direction as the original

numerical solver - i.e. moving forward in time - or, instead, they solve a new system that goes backwards in time. Figure 1 displays a classification of some methods under this two-fold classification. In the following section, we are going to explore more in detail these methods.

It is important to note that if all the methods we explore in this section are mathematically correct, *that does not imply they are numerically stable*. These statements applied to methods based on pure automatic differentiation as well as adjoint methods. We are going to explore this consideration in more detail in section 4.

3.1 Preliminaries

Consider a system of ordinary differential equations (ODEs) given by

$$\frac{du}{dt} = f(u, \theta, t), \quad (1)$$

where $u \in \mathbb{R}^n$ is the unknown solution; f is a function that depends on the state u , some parameter $\theta \in \mathbb{R}^p$, and an independent variable t which we will refer as time, but it can represent another quantity; and with initial condition $u(t_0) = u_0$. Here n denotes the total number of ODEs and p the size of a parameter embedded in the functional form of the differential equation. Although here we consider the case of ODEs, that is, when the derivatives are just with respect to the time variable t , the ideas presented here can be extended of the case of partial differential equations (for example, via the method of lines¹) and algebraic differential equations (ADE). Except for a minority of functions $f(u, \theta, t)$, solutions of the Equation (1) need to be computed using a numerical solver.

We are interested in computing the gradient of a given function $L(u(\cdot, \theta))$ with respect to the parameter θ . Here we are using the letter L to emphasize that in many cases this will be a loss function, but without loss of generality this includes a broader class of functions.

- **Empirical loss functions.** This is usually a real-valued function that quantifies the accuracy or prediction power of a given model. Examples of loss functions include the squared error

$$L(u(\cdot, \theta)) = \frac{1}{2} \|u(t_1; \theta) - u^{\text{target}(t_1)}\|_2^2, \quad (2)$$

where $u^{\text{target}(t_1)}$ is the desired target observation at some later time t_1 . More generally, we can evaluate the loss function at points of the time series for which we have observations,

$$L(u(\cdot, \theta)) = \frac{1}{2} \sum_{i=1}^N \|u(t_i; \theta) - u^{\text{target}(t_i)}\|_2^2. \quad (3)$$

We can also consider the continuous evaluated loss function of the form

$$L(u(\cdot, \theta)) = \int_{t_0}^{t_1} h(u(t; \theta), \theta) dt, \quad (4)$$

with h being a function that quantifies the contribution of the error term at every time $t \in [t_0, t_1]$. Defining a loss function where just the empirical error is penalized is

known as trajectory matching. Other methods like gradient matching and generalized smoothing the loss depends on smooth approximations of the trajectory and their derivatives.

- **Likelihood profiles.** In statistical models, it is common to assume that observations correspond to noisy observations of the underlying dynamical system, $y_i = u(t_i; \theta) + \varepsilon_i$, with ε_i errors or residual that are independent of each other and of the trajectory $u(\cdot; \theta)$.⁵¹ If $p(y|t, \theta)$ is the probability distribution of y , maximum likelihood estimation consists in finding the parameter θ as

$$\theta^* = \operatorname{argmax}_{\theta} \ell(y|\theta) = \prod_{i=1}^n p(y_i|\theta, t_i). \quad (5)$$

When $\varepsilon \sim N(0, \sigma_i^2)$ is Gaussian, the maximum likelihood principle is the same as minimizing $-\log \ell(y|\theta)$ which results in the mean squared error

$$\theta^* = \operatorname{argmin}_{\theta} \{-\log \ell(y|\theta)\} = \operatorname{argmin}_{\theta} \sum_{i=1}^n (y_i - u(t_i; \theta))^2. \quad (6)$$

Provided with a prior distribution $p(\theta)$ for the parameter θ , we can further compute a posterior distribution for θ given the observations y_1, y_2, \dots, y_n following Bayes theorem

$$p(\theta|y) = \frac{p(y|\theta)p(\theta)}{p(y)}. \quad (7)$$

In practice, the posterior is difficult to evaluate and needs to be approximated using Markov chain Monte Carlo sampling methods.¹⁸

- **Quantity of interest.** Another important example is when L returns the value of the solution at one or many points, which is useful when we want to know how the solution itself changes as we move the parameter values.
- **Diagnosis of the solution.** In many cases we are interested in optimizing the value of some variable that is a function of the solution of a differential equation. This is the case in design control theory, a popular approach in aerodynamics modelling where goals include maximizing the speed of an airplane or the lift of a wing given the solution of the flow equation for a given geometry profile.³⁰

We are interested in computing the gradient of the loss function with respect to the parameter θ , which can be written using the chain rule as

$$\frac{dL}{d\theta} = \frac{dL}{du} \frac{\partial u}{\partial \theta}. \quad (8)$$

Notice that this includes also the case where we want to compute the gradient of L with respect to the initial condition of the system. The first term on the right-hand side is usually

easy to evaluate since it just involves the partial derivative of the scalar loss function with respect to the solution. For example, for the loss function in Equation (2) this is simply

$$\frac{dL}{du} = u - u^{\text{target}(t_1)}. \quad (9)$$

The second term on the right-hand side is more difficult to compute and it is usually referred to as the *sensitivity*,

$$s = \frac{\partial u}{\partial \theta} = \begin{bmatrix} \frac{\partial u_1}{\partial \theta_1} & \cdots & \frac{\partial u_1}{\partial \theta_p} \\ \vdots & \ddots & \vdots \\ \frac{\partial u_n}{\partial \theta_1} & \cdots & \frac{\partial u_n}{\partial \theta_p} \end{bmatrix} \in \mathbb{R}^{n \times p}. \quad (10)$$

Notice here the distinction between the total derivative (indicated with the d) and partial derivative symbols (∂). When a function depends on more than one argument, we are going to use the partial derivative symbol to emphasize this distinction (e.g., Equation (10)). On the other side, when this is not the case, we will use the total derivative symbol (e.g., Equation (9)). Also notice that the sensitivity s defined in Equation (10) is what is called a *Jacobian*, that is, a matrix of first derivatives for general vector-valued functions.

3.2 Finite differences

The simplest way of evaluating a derivative is by computing the difference between the evaluation of the function at a given point and a small perturbation of the function. In the case of a loss function, we can approximate

$$\frac{dL}{d\theta_i}(\theta) \approx \frac{L(\theta + \varepsilon e_i) - L(\theta)}{\varepsilon}, \quad (11)$$

with e_i the i -th canonical vector and ε the stepsize. Even better, it is easy to see that the centered difference scheme

$$\frac{dL}{d\theta_i}(\theta) \approx \frac{L(\theta + \varepsilon e_i) - L(\theta - \varepsilon e_i)}{2\varepsilon}, \quad (12)$$

leads also to a more precise estimation of the derivative. While Equation (11) gives to an error of magnitude $\mathcal{O}(\varepsilon)$, the centered differences schemes improves to $\mathcal{O}(\varepsilon^2)$.²

However, there are a series of problems associated with this approach. The first one is due to how this scales with the number of parameters p . Each directional derivative requires the evaluation of the loss function L twice. For the centered differences approach in Equation (12), this requires a total of $2p$ function evaluations, which at the same time demands solving the differential equation in forward mode each time for a new set of parameters.

A second problem is due to rounding errors. Every computer ultimately stores and manipulates numbers using floating points arithmetic.²¹ Equations (11) and (12) involve the subtraction of two numbers that are very close to each other, which leads to large cancellation errors for small values of ε that are amplified by the division by ε . On the other hand, large values of the stepsize give inaccurate estimations of the gradient. Finding the optimal value of ε that trade-offs these two effects is sometimes called the *stepsize dilemma*.⁴³ Due

to this, some heuristics and algorithms have been introduced in order to pick the value of ε ,^{3,27,43} Some of these methods require some a priori knowledge about the function to be differentiated, and others are based on arbitrary historical rules. If many analytical functions, like polynomials and trigonometric functions, can be computed with machine precision, numerical solutions of differential equations have errors larger than machine precision, which leads to inaccurate estimations of the gradient when ε is too small. We will further emphasize this point in Section 4.

Even with all these caveats, finite differences can be useful when computing Jacobian-vector products. Given a Jacobian matrix $J = \frac{\partial f}{\partial u}$ (or the sensitivity $s = \frac{\partial u}{\partial \theta}$) and a vector v , the product Jv corresponding to the directional derivative can be approximated as

$$Jv \approx \frac{f(u + \varepsilon v, \theta, t) - f(u, \theta, t)}{\varepsilon} \quad (13)$$

This approach is used in numerical solvers based on Krylow methods, where linear systems are solved by iterative solving matrix-vectors products.²⁹

3.3 Complex step differentiation

An alternative to finite differences that avoids rounding errors is based on complex variable analysis. The first proposals originated in 1967 using the Cauchy integral theorem involving the numerical evaluation of a complex-valued integral.^{39,40} A new approach recently emerged that uses the Taylor expansion of a function to define its complex generalization.^{42,61} Assuming that we have one single scalar parameter $\theta \in \mathbb{R}$, then the function $L(\theta)$ can be expanded as the Taylor expansion

$$L(\theta + i\varepsilon) = L(\theta) + i\varepsilon L'(\theta) - \frac{1}{2}L''(\theta)\varepsilon^2 + \mathcal{O}(\varepsilon^3), \quad (14)$$

where i is the imaginary unit satisfying $i^2 = -1$. From this equation, we can observe that many factors vanish when we compute the imaginary part $\text{Im}(L(\theta + i\varepsilon))$, which leads to

$$L'(\theta) = \frac{\text{Im}(L(\theta + i\varepsilon))}{\varepsilon} + \mathcal{O}(\varepsilon^2) \quad (15)$$

The method of *complex step differentiation* consists then in estimating the gradient as $\text{Im}(L(\theta + i\varepsilon))/\varepsilon$ for a small value of ε . Besides the advantage of being a method with precision $\mathcal{O}(\varepsilon^2)$, the complex step method avoids subtracting cancellation error and then the value of ε can be reduced to almost machine precision error without affecting the calculation of the derivative. Extension to higher order derivatives can be done by introducing multicomplex variables.³⁶

3.4 Automatic differentiation

Automatic differentiation (AD) is a technology that allows computing gradients through a computer program.²⁵ The main idea is that every computer program manipulating numbers can be reduced to a sequence of simple algebraic operations that have straightforward

derivative expressions, based upon elementary rules of differentiation. The derivatives of the outputs of the computer program with respect to their inputs are then combined using the chain rule. One advantage of AD systems is that we can automatically differentiate programs that include control flow, such as branching, loops or recursions. This is because at the end of the day, any program can be reduced to a trace of input, intermediate and output variables.⁵

Depending if the concatenation of these gradients is done as we execute the program (from input to output) or in a later instance where we trace-back the calculation from the end (from output to input), we are going to talk about *forward* or *reverse* AD, respectively.

3.4.1 Forward mode

Forward mode AD can be implemented in different ways depending on the data structures we use at the moment of representing a computer program. Examples of these data structures include dual numbers and Wengert lists (see⁵ for a good review on these methods).

Dual numbers

Let us first consider the case of dual numbers. The idea is to extend the definition of a numerical variable that takes a certain value to also carry information about its derivative with respect to certain scalar parameter $\theta \in \mathbb{R}$. We can define an abstract type, defined as a dual number, composed of two elements: a *value* coordinate x_1 that carries the value of the variable and a *derivative* coordinate x_2 with the value of the derivative $\frac{\partial x_1}{\partial \theta}$. Just as complex number, we can represent dual numbers in the vectorial form (x_1, x_2) or in the rectangular form

$$x_\epsilon = x_1 + \epsilon x_2 \quad (16)$$

where ϵ is an abstract number with the properties $\epsilon^2 = 0$ and $\epsilon \neq 0$. This last representation is quite convenient since it naturally allow us to extend algebraic operations, like addition and multiplication, to dual numbers. For example, given two dual numbers $x_\epsilon = x_1 + \epsilon x_2$ and $y_\epsilon = y_1 + \epsilon y_2$, it is easy to derive using the fact $\epsilon^2 = 0$ that

$$x_\epsilon + y_\epsilon = (x_1 + y_1) + \epsilon (x_2 + y_2) \quad x_\epsilon y_\epsilon = x_1 y_1 + \epsilon (x_1 y_2 + x_2 y_1). \quad (17)$$

From these last examples, we can see that the derivative component of the dual number carries the information of the derivatives when combining operations. For example, suppose than in the last example the dual variables x_2 and y_2 carry the value of the derivative of x_1 and x_2 with respect to a parameter θ , respectively.

Intuitively, we can think about ϵ as being a differential in the Taylor expansion:

$$\begin{aligned} f(x_1 + \epsilon x_2) &= f(x_1) + \epsilon x_2 f'(x_1) + \epsilon^2 \cdot (\dots) \\ &= f(x_1) + \epsilon x_2 f'(x_1) \end{aligned} \quad (18)$$

When computing first order derivatives, we can ignore everything of order ϵ^2 or larger, which is represented in the condition $\epsilon^2 = 0$. This implies that we can use dual numbers to

implement forward AD through a numerical algorithm. We will explore how this is carried in Section 4.

Computational graph

An useful way of representing a computer program is via a computational graph with intermediate variables that relate the input and output variables. Most scalar functions of interest can be represented in this factorial form as a acyclic directed graph with nodes associated to variables and edges to atomic operations,^{23,25} known as Kantorovich graph.³³ We can define $v_1, v_2, \dots, v_p = \theta_1, \theta_2, \dots, \theta_p$ the input set of variables; v_{p+1}, \dots, v_{m-1} the set of all the intermediate variables, and finally $v_m = L(\theta)$ the final output of a computer program. This can be done in such a way that the order is strict, meaning that each variable v_i is computed just as a function of the previous variables v_j with $j < i$. Once the graph is constructed, we can compute the derivative of every node with respect to other using Bauer formula^{4,47}

$$\frac{\partial v_j}{\partial v_i} = \sum_{\substack{\text{paths } w_0 \rightarrow w_1 \rightarrow \dots \rightarrow w_K \\ \text{with } w_0 = v_i, w_K = v_j}} \prod_{k=0}^{K-1} \frac{\partial w_{k+1}}{\partial w_k}, \quad (19)$$

where the sum is calculated with respect to all the directed paths in the graph connecting the input and target node. Instead of evaluating the last expression for all possible path, a simplification is to increasingly evaluate $j = p + 1, \dots, m$ using the recursion

$$\frac{\partial v_j}{\partial v_i} = \sum_{w \text{ such that } w \rightarrow v_j} \frac{\partial v_j}{\partial w} \frac{\partial w}{\partial v_i} \quad (20)$$

Since every variable node w such that $w \rightarrow v_j$ is an edge of the computational graph have index less than j , we can iterate this procedure as we run the computer program and solve for both the function and its gradient. This is possible because in forward mode the term $\frac{\partial w}{\partial v_i}$ has been computed in a previous iteration, while $\frac{\partial v_j}{\partial w}$ can be evaluated at the same time the node v_j is computed based on only the value of the parent variable nodes.

3.4.2 Backward mode

Backward mode AD is also known as the adjoint of cotangent linear mode or backpropagation in the field of machine learning. The reverse mode of automatic differentiation has been introduced in different contexts²⁴ and materializes the observation made by Phil Wolfe that if the chain rule is implemented in reverse mode, then the ratio between the computation of the gradient of a function and the function itself can be bounded by a constant that do not depend of the number of parameters to differentiate,^{23,66} a point known as *cheap gradient principle*.²⁴ Given a directional graph of operations defined by a Wengert list,⁶⁵ we can compute gradients of any given function in the same fashion as Equation (20) but in backwards mode as

$$\bar{v}_i = \frac{\partial \ell}{\partial v_i} = \sum_{w: v \rightarrow w \in G} \frac{\partial w}{\partial v} \bar{w}. \quad (21)$$

Here we have introduced the notation $\bar{w} = \frac{\partial \ell}{\partial w}$ to indicate that the partial derivative is always of the final loss function with respect to the different program variables. Since in backwards

AD the values of $\bar{\omega}$ are being updated in reverse order, in order to evaluate the terms $\frac{\partial \omega}{\partial v}$ we need to know the state value of all the argument variables v of ω , which need to be stored in memory during the evaluation of the function in order to be able to apply backward AD.

Another way of implementing backwards AD is by defining a *pullback* function,²⁸ a method also known as *continuation-passing style*.⁶⁴ In the backward step, this executes a series of function calls, one for each elementary operation. If one of the nodes in the graph w is the output of an operation involving the nodes v_1, \dots, v_m , where $v_i \rightarrow w$ are all nodes in the graph, then the pullback $\bar{v}_1, \dots, \bar{v}_m = \mathcal{B}_w(\bar{w})$ is a function that accepts gradients with respect to w (defined as \bar{w}) and returns gradients with respect to each v_i (\bar{v}_i) by applying the chain rule. Consider the example of the multiplicative operation $w = v_1 \times v_2$. Then

$$\bar{v}_1, \bar{v}_2 = v_2 \times \bar{w}, \quad v_1 \times \bar{w} = \mathcal{B}_w(\bar{w}), \quad (22)$$

which is equivalent to using the chain rule as

$$\frac{\partial \ell}{\partial v_1} = \frac{\partial}{\partial v_1}(v_1 \times v_2) \frac{\partial \ell}{\partial w}. \quad (23)$$

3.4.3 AD connection with JVPs and VJPs

When working with unit operations that involve matrix operations dealing with vectors of different dimensions, the order in which we apply the chain rule matters.¹⁹ When computing a gradient using AD, we can encounter vector-Jacobian products (VJPs) or Jacobian-vector products (JVP). As their name indicates, the difference between them regards the fact that the quantity we are interested in computing is described by the product of a Jacobian by a vector on the left side (VJP) or the right (JVP).

For nested functions, the Jacobian is described as the product of multiple Jacobian using the chain rule. In this case, the full gradient is computed as the chain product of vectors and Jacobians. Let us consider for example the case of a loss function $L : \mathbb{R}^p \mapsto \mathbb{R}$ that can be decomposed as $L(\theta) = \ell \circ g_k \circ \dots \circ g_2 \circ g_1(\theta)$, with $\ell : \mathbb{R}^{d_k} \mapsto \mathbb{R}$ the final evaluation of the loss function after we apply in order a sequence of intermediate functions $g_i : \mathbb{R}^{d_{i-1}} \mapsto \mathbb{R}^{d_i}$, $d_0 = p$. Now, using the chain rule, we can calculate the gradient of the final loss function as

$$\nabla_{\theta} L = \nabla \ell \cdot Dg_k \cdot Dg_{k-1} \cdot \dots \cdot Dg_2 \cdot Dg_1, \quad (24)$$

with Dg_i the Jacobians of each nested function. Notice that in the last equation, $\nabla \ell \in \mathbb{R}^{d_k}$ is a vector. In order to compute $\nabla_{\theta} L$, we can solve the multiplication starting from the right side, which will correspond to multiple the Jacobians forward from Dg_1 to Dg_k , or from the left side, moving backwards. The important aspect of this last case is that we will always be computing VJPs, since $\nabla \ell \in \mathbb{R}^{d_k}$ is a vector. Since VJPs are easier to evaluate than full Jacobians, the backward mode will be in general faster (see Figure 2). For general rectangular matrices $A \in \mathbb{R}^{d_1 \times d_2}$ and $B \in \mathbb{R}^{d_2 \times d_3}$, the cost of the matrix multiplication AB is $\mathcal{O}(d_1 d_2 d_3)$ (more efficient algorithms exist but this does not impact these results). This implies that forward AD requires a total of

$$d_2 d_1 n + d_3 d_2 p + \dots + d_k d_{k-1} p + d_k p = \mathcal{O}(p) \quad (25)$$




Figure 2: Comparison between forward and backward AD. Changing the order of how we multiply the Jacobians change the total number of floating-point operations, which leads to different computational complexities between forward and backward mode. However, backwards mode requires storing in memory information about the forward execution of the program, while forward mode can update the gradient on running time.

operations, while backwards mode AD requires

$$d_k d_{k-1} + d_{k-1} d_{k-2} + \dots + d_2 d_1 + d_1 n = \mathcal{O}(1) \quad (26)$$

operations, where the \mathcal{O} is with respect to the variable p .

When the function to differentiate has a larger input space than output, AD in backward mode is more efficient as it propagates the chain rule by computing VJPs, the reason why backwards AD is more used in modern machine learning. However, notice that backwards mode AD requires us to save the solution through the forward run in order to run backwards afterwards,⁶ while in forward mode we can just evaluate the gradient as we iterate our sequence of functions. This problem can be overcome with a good checkpointing scheme, somewhat we will discuss later. This means that for problems with a small number of parameters, forward mode can be faster and more memory-efficient than backwards AD.

3.5 Symbolic differentiation

Sometimes AD is compared against symbolic differentiation. According to,³⁷ these two are the same and the only difference is in the data structures used to implement them, while¹⁴ suggests that AD is symbolic differentiation performed by a compiler.

3.6 Sensitivity equations

An easy way to derive an expression for the sensitivity s is by deriving the sensitivity equations,⁵¹ a method also referred to as continuous local sensitivity analysis (CSA). If we consider the original system of ODEs and we differentiate with respect to θ , we then obtain

$$\frac{d}{d\theta} \frac{du}{dt} = \frac{d}{d\theta} f(u(\theta), \theta, t) = \frac{\partial f}{\partial \theta} + \frac{\partial f}{\partial u} \frac{\partial u}{\partial \theta}, \quad (27)$$

that is

$$\frac{ds}{dt} = \frac{\partial f}{\partial u} s + \frac{\partial f}{\partial \theta}. \quad (28)$$

By solving the sensitivity equation at the same time we solve the original differential equation for $u(t)$, we ensure that by the end of the forward step we have calculated both $u(t)$ and $s(t)$. This also implies that as we solve the model forward, we can ensure the same level of numerical precision for the two of them.

In opposition to the methods previously introduced, the sensitivity equations find the gradient by solving a new set of continuous differential equations. Notice also that the obtained sensitivity $s(t)$ can be evaluated at any given time t . This method can be labeled as forward, since we solve both $u(t)$ and $s(t)$ as we solve the differential equation forward in time, without the need of backtracking any operation through the solver.

For systems of equations with few number of parameters, this method is useful since the system of equations composed by Equations (1) and (28) can be solved in $\mathcal{O}(np)$ using the same precision for both solution and sensitivity numerical evaluation. Furthermore, this method does not required saving the solution in memory, so it can be solved purely in forward mode without backtracking operations. However, notice that the term $\frac{\partial f}{\partial u} s$ is in general difficult to compute.

It is important to remark that the sensitivity equations can be also solved in discrete forward mode by numerically discretizing the original ODE and later deriving the discrete sensitivity equations. For most cases, this leads to the same result that in the continuous case.⁶⁸

3.7 Adjoint methods

For complex and large systems, computing the gradient directly on top of the numerical solver (for example, using AD) can be memory expensive since the large number of function evaluations required by the solver and the later store of the intermediate states. For these cases, the adjoint-based method allows us to compute the gradients of a loss function by instead computing an intermediate variable (the adjoint) that serves as a bridge between the solution of the ODE and the final sensitivity.

There is a large family of adjoint methods that in first order we can classify them between discrete and continuous adjoints. The former usually arises as the numerical discretization of the latter, and in general, these two give different different computational results.⁶⁰ Different results exist regarding the consistency or inconsistency between the two approaches, and this usually depends on the ODE and equation. Proofs of the consistency of discrete adjoint methods for Runge-Kutta methods have been provided in.^{55,56} Depending on the choice of

the Runge-Kutta coefficients, we can have a numerical scheme that is both consistent for the original equation and consistent/inconsistent for the adjoint.²⁶

3.7.1 Discrete adjoint method

Also known as the adjoint state method, it is another example of a discrete method that aims to find the gradient by solving an alternative system of linear equations, known as the *adjoint equations*, at the same time that we solve the original system of linear equations defined by the numerical solver. These methods are extremely popular in optimal control theory in fluid dynamics, for example for the design of geometries for vehicles and airplanes that optimize performance.^{15,20} This approach follows the discretize-optimize approach, meaning that we first discretize the system of continuous ODEs and then solve on top of these linear equations.²⁰

Discrete differential equation

The derivation of the discrete adjoint equations is carried once the numerical scheme for solving Equation (1) has been specified. Given a discrete sequence of timesteps t_0, t_1, \dots, t_N , we evaluate the solution at $u_i = u(t_i; \theta)$. Some of the most common numerical solvers include multistep linear solvers of the form

$$\sum_{i=0}^{K_1} \alpha_{ni} u_{n-i} + h_n \sum_{i=0}^{K_2} \beta_{ni} f(u_{n-i}, \theta, t_{n-i}) = 0. \quad (29)$$

and Runge-Kutta methods with

$$u_{n+1} = u_n + h \sum_{i=1}^s b_i k_i \quad (30)$$

$$k_i = f \left(u_n + \sum_{j=1}^s a_{ij} k_j, \theta, t_n + c_n h \right) \quad i = 1, 2, \dots, s. \quad (31)$$

The former is linear in f , which for example is not the case in Runge-Kutta methods with intermediate evaluations.¹ Explicit methods are characterized by $\beta_{n,0} = 0$ for the multistep and $a_{ij} = 0$ if $i \leq j$ for Runge-Kutta methods, otherwise, the method is implicit.

For multistep methods, solving the differential equation implies to be able to solve the system of constraints

$$g_i(u_i; \theta) = u_i - h\beta_{n0}f(u_i, \theta, t_i) - \alpha_i = 0 \quad (32)$$

where α_i has includes the information of all the past iterations. This system can be solved sequentially, by solving for u_i in increasing order of index using Newton method. If we call the super-vector $U = (u_1, u_2, \dots, u_N) \in \mathbb{R}^{nN}$, we can combine all these equations into one single system of equations $G(U) = (g_1(u_1; \theta), \dots, g_N(u_N; \theta)) = 0$.

In the simplest case where the algebraic set of equations is linear and we can write $g_i(u_{i+1}; \theta) = u_{i+1} - A_i(\theta) u_i - b_i$ with $A_i \in \mathbb{R}^{n \times n}$ and $b_i \in \mathbb{R}^n$ defined by the numerical solver,

the condition $G(U) = 0$ simplifies to the linear system of equations

$$A(\theta)U = \begin{bmatrix} \mathbb{I}_{n \times n} & 0 & & & \\ -A_1 & \mathbb{I}_{n \times n} & 0 & & \\ & -A_2 & \mathbb{I}_{n \times n} & 0 & \\ & & & \ddots & \\ & & & -A_{N-1} & \mathbb{I}_{n \times n} \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ u_3 \\ \vdots \\ u_N \end{bmatrix} = \begin{bmatrix} A_0 u_0 + b_0 \\ b_1 \\ b_2 \\ \vdots \\ b_{N-1} \end{bmatrix} = b(\theta), \quad (33)$$

with $\mathbb{I}_{n \times n}$ the identity matrix of size $n \times n$. It is important to notice that in most cases, the matrix $A(\theta)$ is quite large and mostly sparse. If this representation of the discrete differential equation is quite convenient for mathematical manipulations, at the moment of solving the system we will rely in iterative solvers that save memory and computation.

Adjoint state equations

We are interested in differentiating a function $L(U, \theta)$ with respect to the parameter θ . Since here U is the discrete set of evaluations of the solution, examples of loss functions now include

$$L(U, \theta) = \frac{1}{2} \sum_{i=1}^N \|u_i - u_i^{\text{obs}}\|^2, \quad (34)$$

with u_i^{obs} the observed time-series. We further need to impose the constraint that the solution satisfies the algebraic linear equation $G(U; \theta) = 0$. Now,

$$\frac{dL}{d\theta} = \frac{\partial L}{\partial \theta} + \frac{\partial L}{\partial U} \frac{\partial U}{\partial \theta}, \quad (35)$$

and also for the constraint $G(U; \theta) = 0$ we can derive

$$\frac{dG}{d\theta} = \frac{\partial G}{\partial \theta} + \frac{\partial G}{\partial U} \frac{\partial U}{\partial \theta} = 0 \quad (36)$$

which is equivalent to

$$\frac{\partial U}{\partial \theta} = - \left(\frac{\partial G}{\partial U} \right)^{-1} \frac{\partial G}{\partial \theta}. \quad (37)$$

If we replace this last expression into equation (35), we obtain

$$\frac{dL}{d\theta} = \frac{\partial L}{\partial \theta} - \underbrace{\frac{\partial L}{\partial U}}_{\text{vector}} \left(\frac{\partial G}{\partial U} \right)^{-1} \frac{\partial G}{\partial \theta}. \quad (38)$$

The important trick in the adjoint state methods is to observe that in this last equation, the right-hand side can be resolved as a vector-Jacobian product (VJP). Instead of computing the product of the matrices $\left(\frac{\partial G}{\partial U} \right)^{-1}$ and $\frac{\partial G}{\partial \theta}$, it is computationally more efficient first to compute the resulting vector from the operation $\frac{\partial L}{\partial U} \left(\frac{\partial G}{\partial U} \right)^{-1}$ and then multiply this by $\frac{\partial G}{\partial \theta}$. This is what leads to the definition of the adjoint $\lambda \in \mathbb{R}^{nN}$ as the solution of the linear system of equations

$$\left(\frac{\partial G}{\partial U} \right)^T \lambda = \left(\frac{\partial L}{\partial U} \right)^T, \quad (39)$$

that is,

$$\lambda^T = \frac{\partial L}{\partial U} \left(\frac{\partial G}{\partial U} \right)^{-1}. \quad (40)$$

Finally, if we replace Equation (40) into (38), we obtain

$$\frac{dL}{d\theta} = \frac{\partial L}{\partial \theta} - \lambda^T \frac{\partial G}{\partial \theta}. \quad (41)$$

The important trick to notice here is the rearrangement of the multiplicative terms involved in equation (38). Computing the full Jacobian/sensitivity $\partial u / \partial \theta$ will be computationally expensive and involves the product of two matrices. However, we are not interested in the calculation of the Jacobian, but instead in the VJP given by $\frac{\partial L}{\partial U} \frac{\partial U}{\partial \theta}$. By rearranging these terms, we can make the same computation more efficient.

Notice that the algebraic equation of the adjoint λ in Equation (39) is a linear system of equations even when the original system $G(U) = 0$ was not necessarily linear in U . For the linear system of discrete equations $G(U; \theta) = A(\theta)U - b(\theta) = 0$, we have³²

$$\frac{\partial G}{\partial \theta} = \frac{\partial A}{\partial \theta} U - \frac{\partial b}{\partial \theta}, \quad (42)$$

so the desired gradient in Equation (41) can be computed as

$$\frac{dL}{d\theta} = \frac{\partial L}{\partial \theta} - \lambda^T \left(\frac{\partial A}{\partial \theta} U - \frac{\partial b}{\partial \theta} \right) \quad (43)$$

with λ the solution of the linear system (Equation (39))

$$A(\theta)^T \lambda = \begin{bmatrix} \mathbb{I}_{n \times n} & -A_1^T & & & \\ 0 & \mathbb{I}_{n \times n} & -A_2^T & & \\ & 0 & \mathbb{I}_{n \times n} & -A_3^T & \\ & & & \ddots & -A_{N-1}^T \\ & & & 0 & \mathbb{I}_{n \times n} \end{bmatrix} \begin{bmatrix} \lambda_1 \\ \lambda_2 \\ \lambda_3 \\ \vdots \\ \lambda_N \end{bmatrix} = \begin{bmatrix} u_1 - u_1^{\text{obs}} \\ u_2 - u_2^{\text{obs}} \\ u_3 - u_3^{\text{obs}} \\ \vdots \\ u_N - u_N^{\text{obs}} \end{bmatrix} = \frac{\partial L}{\partial U}. \quad (44)$$

This is a linear system of equations with the same size of the original $A(\theta)U = b(\theta)$, but involving the adjoint matrix A^T . Computationally this also means that if we can solve the original system of discretized equations then we can also solve the adjoint. One way of doing this is relying on matrix factorization. Using the LU factorization we can write the matrix $A(\theta)$ as the product of a lower and upper triangular matrices $A(\theta) = LU$, which then can be also used for solving the adjoint equation since $A^T(\theta) = U^T L^T$. Another more natural way of finding the adjoints λ is by noticing that the system of equations (44) is equivalent to the iterative scheme

$$\lambda_i = A_i^T \lambda_{i+1} + (u_i - u_i^{\text{obs}}) \quad (45)$$

with initial condition λ_N . This means that we can solve the adjoint equation in backwards mode, starting from the final state λ_N and computing the values of λ_i in decreasing index order. Notice that this procedure requires to know the value of u_i at any given timestep.

3.7.2 Continuous adjoint method

The continuous adjoint method, also known as continuous adjoint sensitivity analysis (CASA), operates by defining a convenient set of new differential equations for the adjoint variable and using this to compute the gradient in a more efficient manner. Mathematically speaking, the adjoint equations can be derived from a duality or Lagrangian point of view.²⁰ We prefer to derive the equation using the former methods since we believe it gives better insights to how the method works and also allow to generalize to other user cases. The derivation of both the discrete and continuous adjoint methods using Lagrangian multipliers can be found in Appendix A. We encourage the interested reader to make the effort of following how the continuous adjoint method follows the same logic than the discrete methods, but where the discretization of the differential equation does not happen until the very last step, when the solutions of the differential equations involved need to be numerically evaluated.

Consider an integrated loss function of the form

$$L(u; \theta) = \int_{t_0}^{t_1} h(u(t; \theta), \theta) dt \quad (46)$$

and its derivative with respect to the parameter θ given by the following integral involving the sensitivity matrix $s(t)$:

$$\frac{dL}{d\theta} = \int_{t_0}^{t_1} \left(\frac{\partial h}{\partial \theta} + \frac{\partial h}{\partial u} s(t) \right) dt. \quad (47)$$

Just as in the case of the discrete adjoint method, the complicated term to evaluate in the last expression is the sensitivity (Equation (10)). Just as in the case of the discrete adjoint method, the trick is to evaluate the VJP $\frac{\partial h}{\partial u} \frac{\partial u}{\partial \theta}$ by defining an intermediate adjoint variable. The continuous adjoint equation now is obtained by finding the dual/adjoint equation of the sensitivity equation using the weak formulation of Equation (28). The adjoint equation is obtained by writing the sensitivity equation in the form

$$\int_{t_0}^{t_1} \lambda(t)^T \left(\frac{ds}{dt} - f(u, \theta, t) s - \frac{\partial f}{\partial \theta} \right) dt = 0, \quad (48)$$

where this equation must be satisfied for every function $\lambda(t)$ in order for Equation (61) to be true. The next step is to get rid of all time derivative applied to the sensitivity $s(t)$ using integration by parts:

$$\int_{t_0}^{t_1} \lambda(t)^T \frac{ds}{dt} dt = \lambda(t_1)^T s(t_1) - \lambda(t_0)^T s(t_0) - \int_{t_0}^{t_1} \frac{d\lambda^T}{dt} s(t) dt. \quad (49)$$

Replacing this last expression into Equation (48) we obtain

$$\int_{t_0}^{t_1} \left(-\frac{d\lambda^T}{dt} - \lambda(t)^T f(u, \theta, t) \right) s(t) dt = \int_{t_0}^{t_1} \lambda(t)^T \frac{\partial f}{\partial \theta} dt - \lambda(t_1)^T s(t_1) + \lambda(t_0)^T s(t_0). \quad (50)$$

At first glance, there is nothing particularly interesting about this last equation. However, both Equations (47) and (50) involved a VJP with $s(t)$. Since Equation (50) must hold for

every function $\lambda(t)$, we can pick $\lambda(t)$ to make the terms involving $s(t)$ in Equations (47) and (50) to perfectly coincide. This is done by defining the adjoint $\lambda(t)$ to be the solution of the new system of differential equations

$$\frac{d\lambda}{dt} = -f(u, \theta, t)^T \lambda - \frac{\partial h^T}{\partial u} \quad \lambda(t_1) = 0. \quad (51)$$

Notice that the adjoint equation is defined with the final condition at t_1 , meaning that it needs to be solved backwards in time. The definition of the adjoint $\lambda(t)$ as the solution of this last ODE simplifies Equation (50) to

$$\int_{t_0}^{t_1} \frac{\partial h}{\partial u} s(t) dt = \lambda(t_0)^T s(t_0) + \int_{t_0}^{t_1} \lambda(t)^T \frac{\partial f}{\partial \theta} dt. \quad (52)$$

Finally, replacing this inside the expression for the gradient of the loss function we have

$$\frac{dL}{d\theta} = \lambda(t_0)^T s(t_0) + \int_{t_0}^{t_1} \left(\frac{\partial h}{\partial \theta} + \lambda^T \frac{\partial f}{\partial \theta} \right) dt \quad (53)$$

The full algorithm to compute the full gradient $\frac{dL}{d\theta}$ can be described as follows:

1. Solve the original differential equation $\frac{du}{dt} = f(u, t, \theta)$;
2. Solve the backwards adjoint differential equation (51);
3. Compute the gradient using Equation (53).

4 Computational implementation

In this section, we are going to address how these different methods are computationally implemented and how to decide which method to use depending on the scientific task. In order to address this point, it is important to make one further distinction of the methods introduced in Section 3 between those that apply direct differentiation at the algorithmic level or those that are based on numerical solvers. The first is easier to implement since they are agnostic with respect to the details of the ODE and its numerical solution; however, they tend to be either inaccurate, memory-expensive, or unfeasible for large models. The family of methods that are based on numerical solvers include the sensitivity equations and the adjoint methods, both discrete and continuous; they are more difficult to implement and for real case applications require complex software implementations, but they are also more accurate and adequate.

4.1 Direct methods

Direct methods are implemented independent of the structure of the ODE and the numerical solver used to solve it.

4.1.1 Finite differences

Finite differences are easy to implement manually, do not require much software support, and provide a direct way of approximating a gradient. In Julia, these methods are implemented in `FiniteDiff.jl` and `FiniteDifferences.jl` and it is recommended to use established libraries than implementing it yourself, since these already include subroutines to determine step-sizes. However, finite differences are less accurate and as costly as forward AD²³ and complex-step differentiation. Figure 3 illustrates the error in computing the gradient of a simple loss function for both true analytical solution and numerical solution of a system of ODEs as a function of the stepsize ε using finite differences and complex-step differentiation.




Figure 3: Absolute relative error when computing the gradient of the function $u(t) = \sin(\omega t)/\omega$ with respect to ω at $t = 10.0$ as a function of the stepsize ε . Here $u(t)$ corresponds to the solution of the differential equation $u'' + \omega^2 u = 0$ with initial condition $u(0) = 0$ and $u'(0) = 1$. The blue dots correspond to the case where this is computed with finite differences. The red and orange lines are for the case where $u(t)$ is numerically computed using the default Tsitouras solver⁶³ from `OrdinaryDiffEq.jl` using different tolerances. The error when using a numerical solver is larger and it is dependent on the numerical precision of the numerical solver.

4.1.2 Forward-mode AD

Implementing forward AD using dual numbers is usually carried out using *operator overloading*.⁴⁵ This means expanding the object associated to a numerical value to include its dual components (derivative) and extending the definition of atomic algebraic functions. In Julia, this can be done by relying on multiple dispatch. The following example illustrates how to define a dual number and its associated binary addition and multiplication extensions.

```

using Base: @kwdef

@kwdef struct DualNumber{F <: AbstractFloat}
    value::F
    derivative::F
end

# Binary sum
Base.:(+)(a::DualNumber, b::DualNumber) = DualNumber(value = a.value + b.value,
    derivative = a.derivative + b.derivative)

# Binary product
Base.:*(a::DualNumber, b::DualNumber) = DualNumber(value = a.value * b.value,
    derivative = a.value*b.derivative + a.derivative*b.value)

```

We can also extend the definition of standard functions by simply applying the chain rule and storing the derivative in the dual variable following Equation (18):

```

function Base.:(sin)(a::DualNumber)
    value = sin(a.value)
    derivative = a.derivative * cos(a.value)
    return DualNumber(value=value, derivative=derivative)
end

```

In the Julia ecosystem, ForwardDiff.jl implements forward mode AD with multidimensional dual numbers.⁵³ Notice that a major limitation of the dual number approach is that we need a dual variable for each variable we want to differentiate. Implementations of forward AD using dual numbers and computational graphs require a number of operations that increases with the number of variables to differentiate, since each computed quantity is accompanied by the corresponding gradient calculations.²³ This consideration also applies to the other forward methods, including finite differences and complex-step differentiation, which makes forward models inefficient when differentiating with respect to many parameters.

Notice that both AD based in dual number and complex-step differentiation introduce an abstract unit (ϵ and i , respectively) associated with the imaginary part of the extender value that carries forward the numerical value of the gradient. Although these methods seem similar, it is important to remark that AD gives the exact gradient, while complex step differentiation relies on numerical approximations that are valid just when the stepsize ϵ is small. The next example shows how the calculation of the gradient of $\sin(x^2)$ is performed by these two methods:

Operation	AD with Dual Numbers	Complex Step Differentiation
x	$x + \epsilon$	$x + i\epsilon$
x^2	$x^2 + \epsilon(2x)$	$x^2 - \epsilon^2 + 2i\epsilon x$
$\sin(x^2)$	$\sin(x^2) + \epsilon \cos(x^2)(2x)$	$\sin(x^2 - \epsilon^2) \cosh(2i\epsilon) + i \cos(x^2 - \epsilon^2) \sinh(2i\epsilon)$

(54)

While the second component of the dual number has the exact derivative of $\sin(x^2)$, it is not until we take $\epsilon \rightarrow 0$ than we obtain the derivative in the imaginary component for the

complex step method

$$\lim_{\varepsilon \rightarrow 0} \frac{1}{\varepsilon} \cos(x^2 - \varepsilon^2) \sinh(2i\varepsilon) = \cos(x^2)(2x). \quad (55)$$

The stepsize dependence of the complex step differentiation method makes it resemble more to finite differences than AD with dual numbers. This difference between the methods also makes the complex step method sometimes more efficient than both finite differences and AD,³⁶ an effect that can be counterbalanced by the number of extra unnecessary operations that complex arithmetic requires (see last column of (54)).⁴² It is also important to remark that many modern software already have support for complex number arithmetic, making complex step differentiation very easy to implement.

4.1.3 Reverse-mode AD

The libraries `ReverseDiff.jl` and `Zygote.jl` use callbacks to compute gradients. When gradients are being computed with less than ~ 100 parameters, the former is faster (see documentation).

5

4.2 Solver-based methods

Sensitivity methods based on numerical solvers tend to be better adapted to the structure and properties of the underlying ODE (stiffness, stability, accuracy) but also more difficult to implement. This difficulty arises from the fact that the sensitivity method needs to deal with some numerical and computational considerations, including how to handle matrix/Jacobian-vector products; numerical stability of the forward/backward solver; and memory-time tradeoff. These factors are further exacerbated by the number of ODEs and parameters in the model. Just a few modern scientific software have the capabilities of handling ODE solvers and computing their sensitivities at the same time. They include CVODES within SUNDIALS in C;^{27,58} ODESSA³⁸ and FATODE (discrete adjoints)⁶⁸ both in Fortran; `SciMLSensitivity.jl` in Julia;⁴⁹ `Dolfin-adjoint` based on the FEniCS Project.^{16,44}

It is important to remark that the underlying machinery of all solvers relies on solvers for linear systems of equations, which can be solved in dense, band (sparse), and Krylow mode. Another important consideration is that all these methods have subroutines to compute the VJPs involved in the sensitivity and adjoint equations. This calculation is carried out by another sensitivity method (finite differences, AD) and this also plays a central role at the moment of analyzing the accuracy and stability of the adjoint method.

4.2.1 Sensitivity equation

4.2.2 Solving the adjoint

An equally important consideration when working with adjoints is when these are numerically stable. Some works have shown that continuous adjoints can lead to unstable sensitivities.³¹

Implicit forward schemes can give rise to explicit backwards schemes, leading to unstable solutions for the gradient.

Solving the backwards mode

The bottleneck of this method is the calculation of the adjoint since in order to solve the adjoint equation we need to know $u(t)$ at any given time. Effectively, notice that the adjoint equation involves the terms $f(u, \theta, t)$ and $\frac{\partial h}{\partial u}$ which are both functions of $u(t)$. There are different ways of addressing the evaluation of $u(t)$ during the backwards step.

- (i) **Dense Store.** During the forward model, we can just store in memory all the intermediate states of the numerical solution. This leads to heavy-memory expensive algorithms.
- (ii) **Re-solve.** Solve again the original ODE together with the adjoint as the solution of the reversed augmented system¹¹

$$\frac{d}{dt} \begin{bmatrix} u \\ \lambda \\ \frac{dL}{d\theta} \end{bmatrix} = \begin{bmatrix} -f \\ -\frac{\partial f}{\partial u}^T \lambda - \frac{\partial h}{\partial u}^T \\ -\lambda^T \frac{\partial f}{\partial \theta} - \frac{\partial h}{\partial \theta} \end{bmatrix} \quad \begin{bmatrix} u \\ \lambda \\ \frac{dL}{d\theta} \end{bmatrix} (t_1) = \begin{bmatrix} u(t_1) \\ \frac{\partial L}{\partial u(t_1)} \\ \lambda(t_0)^T s(t_0) \end{bmatrix}. \quad (56)$$

However, computing the ODE backwards can be unstable and lead to large numerical errors.^{35,69}

- (iii) **Checkpointing.** Also known as windowing, checkpointing is a technique that trades off memory and time by saving intermediate states of the solution in the forward pass and recalculating the solution between intermediate states in the backwards mode.^{25,57} This is implemented in `Checkpointing.jl`.⁵⁷

One way of solving this system of equations that ensures stability is by using implicit methods. However, this requires cubic time in the total number of ordinary differential equations, leading to a total complexity of $\mathcal{O}((n+p)^3)$ for the adjoint method. Two alternatives are proposed in,³⁵ the first one called *Quadrature Adjoint* produces a high order interpolation of the solution $u(t)$ as we move forward, then solve for λ backwards using an implicit solver and finally integrating $\frac{dL}{d\theta}$ in a forward step. This reduces the complexity to $\mathcal{O}(n^3 + p)$, where the cubic cost in the number of ODEs comes from the fact that we still need to solve the original stiff differential equation in the forward step. A second but similar approach is to use an implicit-explicit (IMEX) solver, where we use the implicit part for the original equation and the explicit for the adjoint. This method also will have complexity $\mathcal{O}(n^3 + p)$.

Solving the quadrature

Another computational challenge in the computation of the continuous adjoint is how the integral in Equation (53) is numerically evaluated. Some methods save computation by noticing that the last step in the continuous adjoint method of evaluating $\frac{dL}{d\theta}$ is an integral instead of an ODE, and then can be evaluated as such without the need to include it in the tolerance calculation inside the numerical solver.³⁴

Numerical solutions of the integral

$$\int_{t_0}^{t_1} \approx \Sigma \tag{57}$$

Gaussian quadrature is the faster method to evaluate one-dimensional integrals.⁴⁶

Weights and knots are obtained in order to maximize the order in which polynomials are exactly integrated.⁶²

4.2.3 Further considerations

Appendices

A Lagrangian derivation of adjoints

In this section we are going to derive the adjoint equation for both discrete and continuous methods using the Lagrange multiplier trick. Conceptually, the method is the same in both discrete and continuous case, with the difference that we manipulate linear algebra objects for the former and continuous operators for the later.

For the continuous adjoint method, we proceed the same way by writing a new loss function $I(\theta)$ identical to $L(\theta)$ as

$$I(\theta) = L(\theta) - \int_{t_0}^{t_1} \lambda(t)^T \left(\frac{du}{dt} - f(u, \theta, t) \right) dt \quad (58)$$

where $\lambda(t) \in \mathbb{R}^n$ is the Lagrange multiplier of the continuous constraint defined by the differential equation. Now,

$$\frac{dL}{d\theta} = \frac{dI}{d\theta} = \int_{t_0}^{t_1} \left(\frac{\partial h}{\partial \theta} + \frac{\partial h}{\partial u} \frac{\partial u}{\partial \theta} \right) dt - \int_{t_0}^{t_1} \lambda(t)^T \left(\frac{d}{dt} \frac{du}{d\theta} - \frac{\partial f}{\partial u} \frac{du}{d\theta} - \frac{\partial f}{\partial \theta} \right) dt. \quad (59)$$

Notice that the term involved in the second integral is the same we found when deriving the sensitivity equations. We can derive an easier expression for the last term using integration by parts. Using our usual definition of the sensitivity $s = \frac{du}{d\theta}$, and performing integration by parts in the term $\lambda^T \frac{d}{dt} \frac{du}{d\theta}$ we derive

$$\begin{aligned} \frac{dL}{d\theta} = \int_{t_0}^{t_1} \left(\frac{\partial h}{\partial \theta} + \lambda^T \frac{\partial f}{\partial \theta} \right) dt - \int_{t_0}^{t_1} \left(-\frac{d\lambda^T}{dt} - \lambda^T \frac{\partial f}{\partial u} - \frac{\partial h}{\partial u} \right) s(t) dt \\ - \left(\lambda(t_1)^T s(t_1) - \lambda(t_0)^T s(t_0) \right). \end{aligned} \quad (60)$$

Now, we can force some of the terms in the last equation to be zero by solving the following adjoint differential equation for $\lambda(t)^T$ in backwards mode

$$\frac{d\lambda}{d\theta} = - \left(\frac{\partial f}{\partial u} \right)^T \lambda - \left(\frac{\partial h}{\partial u} \right)^T, \quad (61)$$

with final condition $\lambda(t_1) = 0$.

B Supplementaty code

References

1. URI M ASCHER. *Numerical methods for evolutionary differential equations*. SIAM, 2008.
2. URI M. ASCHER and CHEN GREIF. *A First Course in Numerical Methods*. Philadelphia, PA: Society for Industrial and Applied Mathematics, 2011. DOI: 10.1137/9780898719987.
3. RUSSELL R. BARTON. “Computing Forward Difference Derivatives In Engineering Optimization”. In: *Engineering Optimization* 20.3 (1992), pp. 205–224. ISSN: 0305-215X. DOI: 10.1080/03052159208941281.
4. F L BAUER. “Computational Graphs and Rounding Error”. In: *SIAM Journal on Numerical Analysis* 11.1 (1974), pp. 87–96. ISSN: 0036-1429. DOI: 10.1137/0711010.
5. ATILIM GUNES BAYDIN et al. “Automatic differentiation in machine learning: a survey”. In: *arXiv* (2015). DOI: 10.48550/arxiv.1502.05767.
6. C H BENNETT. “Logical Reversibility of Computation”. In: *IBM Journal of Research and Development* 17.6 (1973), pp. 525–532. ISSN: 0018-8646. DOI: 10.1147/rd.176.0525.
7. JEFF BEZANSON et al. “Julia: A Fast Dynamic Language for Technical Computing”. In: *arXiv* (2012). DOI: 10.48550/arxiv.1209.5145.
8. JEFF BEZANSON et al. “Julia: A Fresh Approach to Numerical Computing”. In: *SIAM Review* 59.1 (2017), pp. 65–98. ISSN: 0036-1445. DOI: 10.1137/141000671.
9. JORDI BOLIBAR et al. “Universal Differential Equations for glacier ice flow modelling”. In: (June 2023). DOI: 10.5194/gmd-2023-120. URL: <https://gmd.copernicus.org/preprints/gmd-2023-120/> (visited on 07/11/2023).
10. ANDREW M BRADLEY. *PDE-constrained optimization and the adjoint method*. Tech. rep. Technical Report. Stanford University. <https://cs.stanford.edu/~ambrad...>, 2013.
11. RICKY TQ CHEN et al. “Neural ordinary differential equations”. In: *Advances in neural information processing systems* 31 (2018).
12. K.M. CUFFEY and W.S.B. PATERSON. *The Physics of Glaciers*. Elsevier Science, 2010. ISBN: 978-0-08-091912-6. URL: <https://books.google.fr/books?id=Jca2v1u1EKEC>.
13. RAJ DANDEKAR, CHRIS RACKAUCKAS, and GEORGE BARBASTATHIS. “A Machine Learning-Aided Global Diagnostic and Comparative Tool to Assess Effect of Quarantine Control in COVID-19 Spread”. In: *Patterns* 1.9 (2020), p. 100145. ISSN: 2666-3899. DOI: 10.1016/j.patter.2020.100145.
14. CONAL ELLIOTT. “The simple essence of automatic differentiation”. In: *Proceedings of the ACM on Programming Languages* 2.ICFP (2018), p. 70. DOI: 10.1145/3236765.

15. JONATHAN ELLIOTT and JAIME PERAIRE. “Aerodynamic design using unstructured meshes”. In: *Fluid Dynamics Conference* (1996). This has an example of the hardcore adjoint method implemented for aerodynamics. It may help to read this to see how the adjoint equations is being solved and the size of the problem. DOI: 10.2514/6.1996-1941.
16. P. E. FARRELL et al. “Automated Derivation of the Adjoint of High-Level Transient Finite Element Programs”. In: *SIAM Journal on Scientific Computing* 35.4 (2013), pp. C369–C393. DOI: 10.1137/120873558. eprint: <https://doi.org/10.1137/120873558>. URL: <https://doi.org/10.1137/120873558>.
17. M. GELBRECHT et al. “Differentiable programming for Earth system modeling”. In: *Geoscientific Model Development* 16.11 (2023), pp. 3123–3135. DOI: 10.5194/gmd-16-3123-2023. URL: <https://gmd.copernicus.org/articles/16/3123/2023/>.
18. ANDREW GELMAN et al. *Bayesian data analysis*. CRC press, 2013.
19. RALF GIERING and THOMAS KAMINSKI. “Recipes for adjoint code construction”. In: *ACM Transactions on Mathematical Software (TOMS)* 24.4 (1998), pp. 437–474. ISSN: 0098-3500. DOI: 10.1145/293686.293695.
20. MICHAEL B. GILES and NILES A. PIERCE. “An Introduction to the Adjoint Approach to Design”. In: *Flow, Turbulence and Combustion* 65.3–4 (2000), pp. 393–415. ISSN: 1386-6184. DOI: 10.1023/a:1011430410075.
21. DAVID GOLDBERG. “What every computer scientist should know about floating-point arithmetic”. In: *ACM Computing Surveys (CSUR)* 23.1 (1991), pp. 5–48. ISSN: 0360-0300. DOI: 10.1145/103162.103163.
22. A.N. GORBAN and D.C. WUNSCH. “The general approximation theorem”. In: *1998 IEEE International Joint Conference on Neural Networks Proceedings. IEEE World Congress on Computational Intelligence (Cat. No.98CH36227)*. Vol. 2. 1998, 1271–1274 vol.2. DOI: 10.1109/IJCNN.1998.685957.
23. ANDREAS GRIEWANK. “On Automatic Differentiation”. In: (Feb. 1997).
24. ANDREAS GRIEWANK. “Who invented the reverse mode of differentiation”. In: *Documenta Mathematica, Extra Volume ISMP* 389400 (2012).
25. ANDREAS GRIEWANK and ANDREA WALTHER. *Evaluating derivatives: principles and techniques of algorithmic differentiation*. SIAM, 2008.
26. WILLIAM W. HAGER. “Runge-Kutta methods in optimal control and the transformed adjoint system”. In: *Numerische Mathematik* 87.2 (2000), pp. 247–282. ISSN: 0029-599X. DOI: 10.1007/s002110000178.
27. ALAN C HINDMARSH et al. “SUNDIALS: Suite of nonlinear and differential/algebraic equation solvers”. In: *ACM Transactions on Mathematical Software (TOMS)* 31.3 (2005), pp. 363–396.
28. MICHAEL INNES. “Don’t Unroll Adjoint: Differentiating SSA-Form Programs”. In: *arXiv* (2018).

29. ILSE C. F. IPSEN and CARL D. MEYER. "The Idea Behind Krylov Methods". In: *The American Mathematical Monthly* 105.10 (1998), pp. 889–899. ISSN: 0002-9890. DOI: 10.1080/00029890.1998.12004985.
30. ANTONY JAMESON. "Aerodynamic design via control theory". In: *Journal of Scientific Computing* 3.3 (1988), pp. 233–260. ISSN: 0885-7474. DOI: 10.1007/bf01061285.
31. JAKOB S. JENSEN, PRAVEEN B. NAKSHATRALA, and DANIEL A. TORTORELLI. "On the consistency of adjoint sensitivity analysis for structural optimization of linear dynamic problems". In: *Structural and Multidisciplinary Optimization* 49.5 (2014), pp. 831–837. ISSN: 1615-147X. DOI: 10.1007/s00158-013-1024-4.
32. STEVEN G. JOHNSON. "Notes on Adjoint Methods for 18.335". In: 2012.
33. LEONID VITALEVICH KANTOROVICH. "On a mathematical symbolism convenient for performing machine calculations". In: *Dokl. Akad. Nauk SSSR*. Vol. 113. 4. 1957, pp. 738–741.
34. PATRICK KIDGER, RICKY T. Q. CHEN, and TERRY J LYONS. "'Hey, that's not an ODE': Faster ODE Adjoints via Seminorms". In: *Proceedings of the 38th International Conference on Machine Learning*. Ed. by MARINA MEILA and TONG ZHANG. Vol. 139. Proceedings of Machine Learning Research. PMLR, 18–24 Jul 2021, pp. 5443–5452. URL: <https://proceedings.mlr.press/v139/kidger21a.html>.
35. SUYONG KIM et al. "Stiff neural ordinary differential equations". en. In: *Chaos: An Interdisciplinary Journal of Nonlinear Science* 31.9 (Sept. 2021), p. 093122. ISSN: 1054-1500, 1089-7682. DOI: 10.1063/5.0060697. URL: <https://aip.scitation.org/doi/10.1063/5.0060697> (visited on 02/25/2022).
36. GREGORY LANTOINE, RYAN P. RUSSELL, and THIERRY DARGENT. "Using Multi-complex Variables for Automatic Computation of High-Order Derivatives". In: *ACM Transactions on Mathematical Software (TOMS)* 38.3 (2012), p. 16. ISSN: 0098-3500. DOI: 10.1145/2168773.2168774.
37. SOEREN LAUE. *On the Equivalence of Forward Mode Automatic Differentiation and Symbolic Differentiation*. 2019. DOI: 10.48550/ARXIV.1904.02990. URL: <https://arxiv.org/abs/1904.02990>.
38. JORGE R. LEIS and MARK A. KRAMER. "Algorithm 658: ODESSA—an Ordinary Differential Equation Solver with Explicit Simultaneous Sensitivity Analysis". In: *ACM Trans. Math. Softw.* 14.1 (Mar. 1988), pp. 61–67. ISSN: 0098-3500. DOI: 10.1145/42288.214371. URL: <https://doi.org/10.1145/42288.214371>.
39. J N LYNESS. "Numerical algorithms based on the theory of complex variable". In: *Proceedings of the 1967 22nd national conference on -* (1967), pp. 125–133. DOI: 10.1145/800196.805983.
40. J N LYNESS and C B MOLER. "Numerical Differentiation of Analytic Functions". In: *SIAM Journal on Numerical Analysis* 4.2 (1967), pp. 202–210. ISSN: 0036-1429. DOI: 10.1137/0704019.

41. YINGBO MA et al. “A Comparison of Automatic Differentiation and Continuous Sensitivity Analysis for Derivatives of Differential Equation Solutions”. In: *arXiv:1812.01892 [cs]* (July 2021). arXiv: 1812.01892. URL: <http://arxiv.org/abs/1812.01892> (visited on 02/25/2022).
42. JOAQUIM R. R. A. MARTINS, PETER STURDZA, and JUAN J. ALONSO. “The complex-step derivative approximation”. In: *ACM Transactions on Mathematical Software (TOMS)* 29 (2003), pp. 245–262. ISSN: 0098-3500. DOI: 10.1145/838250.838251.
43. RAVISHANKAR MATHUR. “An analytical approach to computing step sizes for finite-difference derivatives”. PhD thesis. 2012.
44. SEBASTIAN K. MITUSCH, SIMON W. FUNKE, and JØRGEN S. DOKKEN. “dolphin-adjoint 2018.1: automated adjoints for FEniCS and Firedrake”. In: *Journal of Open Source Software* 4.38 (2019), p. 1292. DOI: 10.21105/joss.01292. URL: <https://doi.org/10.21105/joss.01292>.
45. MARTIN NEUENHOFEN. “Review of theory and implementation of hyper-dual numbers for first and second order automatic differentiation”. In: *arXiv* (2018). DOI: 10.48550/arxiv.1801.03614.
46. ALEXANDER NORCLIFFE and MARC PETER DEISENROTH. “Faster Training of Neural ODEs Using Gauß-Legendre Quadrature”. In: *arXiv* (2023). DOI: 10.48550/arxiv.2308.10644.
47. DENIZ OKTAY et al. “Randomized Automatic Differentiation”. In: *arXiv* (2020). DOI: 10.48550/arxiv.2007.10412.
48. DEREK ONKEN and LARS RUTHOTTO. “Discretize-Optimize vs. Optimize-Discretize for Time-Series Regression and Continuous Normalizing Flows”. In: *arXiv* (2020). DOI: 10.48550/arxiv.2005.13420.
49. CHRISTOPHER RACKAUCKAS et al. “Universal differential equations for scientific machine learning”. In: *arXiv preprint arXiv:2001.04385* (2020).
50. M. RAISSI, P. PERDIKARIS, and G.E. KARNIADAKIS. “Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations”. In: *Journal of Computational Physics* 378 (2019), pp. 686–707. ISSN: 0021-9991. DOI: 10.1016/j.jcp.2018.10.045.
51. JAMES RAMSAY and GILES HOOKER. *Dynamic data analysis*. Springer, 2017.
52. BHARATH RAMSUNDAR, DILIP KRISHNAMURTHY, and VENKATASUBRAMANIAN VISWANATHAN. “Differentiable Physics: A Position Piece”. In: *arXiv* (2021). DOI: 10.48550/arxiv.2109.07573.
53. J. REVELS, M. LUBIN, and T. PAPAMARKOU. “Forward-Mode Automatic Differentiation in Julia”. In: *arXiv:1607.07892 [cs.MS]* (2016). URL: <https://arxiv.org/abs/1607.07892>.
54. SEBASTIAN RUDER. “An overview of gradient descent optimization algorithms”. In: *arXiv preprint arXiv:1609.04747* (2016).

55. ADRIAN SANDU. “On the properties of Runge-Kutta discrete adjoints”. In: *Computational Science–ICCS 2006: 6th International Conference, Reading, UK, May 28–31, 2006, Proceedings, Part IV* 6. Springer. 2006, pp. 550–557.
56. ADRIAN SANDU. “Solution of inverse problems using discrete ODE adjoints”. In: *Large-Scale Inverse Problems and Quantification of Uncertainty* (2011), pp. 345–365.
57. MICHEL SCHANEN et al. “Transparent Checkpointing for Automatic Differentiation of Program Loops Through Expression Transformations”. In: (2023). Ed. by JIŘÍ MIKYŠKA et al., pp. 483–497.
58. RADU SERBAN and ALAN C HINDMARSH. “CVODES: the sensitivity-enabled ODE solver in SUNDIALS”. In: *International Design Engineering Technical Conferences and Computers and Information in Engineering Conference*. Vol. 47438. 2005, pp. 257–269.
59. CHAOPENG SHEN et al. “Differentiable modelling to unify machine learning and physical models for geosciences”. In: *Nature Reviews Earth & Environment* (2023), pp. 1–16. DOI: 10.1038/s43017-023-00450-9.
60. ZIV SIRKES and ELI TZIPERMAN. “Finite Difference of Adjoint or Adjoint of Finite Difference?” In: *Monthly Weather Review* 125.12 (1997), pp. 3373–3378. ISSN: 0027-0644. DOI: 10.1175/1520-0493(1997)125<3373:fdoaoa>2.0.co;2.
61. WILLIAM SQUIRE and GEORGE TRAPP. “Using Complex Variables to Estimate Derivatives of Real Functions”. In: 40 (1998), pp. 110–112. ISSN: 0036-1445. DOI: 10.1137/s003614459631241x.
62. JOSEF STOER and ROLAND BULIRSCH. *Introduction to numerical analysis*. Springer, 2002.
63. CH. TSITOURAS. “Runge–Kutta pairs of order 5(4) satisfying only the first column simplifying assumption”. In: *Computers & Mathematics with Applications* 62.2 (2011), pp. 770–775. ISSN: 0898-1221. DOI: 10.1016/j.camwa.2011.06.002.
64. FEI WANG et al. “Backpropagation with Continuation Callbacks: Foundations for Efficient and Expressive Differentiable Programming”. In: *Proceedings of the ACM on Programming Languages* 3.ICFP (2019), p. 96. DOI: 10.1145/3341700.
65. R. E. WENGERT. “A simple automatic derivative evaluation program”. In: *Communications of the ACM* 7.8 (1964), pp. 463–464. ISSN: 0001-0782. DOI: 10.1145/355586.364791.
66. PHILIP WOLFE. “Checking the Calculation of Gradients”. In: *ACM Transactions on Mathematical Software (TOMS)* 8.4 (1982), pp. 337–343. ISSN: 0098-3500. DOI: 10.1145/356012.356013.
67. LENKA ZDEBOROVÁ. “Understanding deep learning is also a job for physicists”. en. In: *Nature Physics* (May 2020). ISSN: 1745-2473, 1745-2481. DOI: 10.1038/s41567-020-0929-2. URL: <http://www.nature.com/articles/s41567-020-0929-2> (visited on 05/29/2020).
68. HONG ZHANG and ADRIAN SANDU. “FATODE: A library for forward, adjoint, and tangent linear integration of ODEs”. In: *SIAM Journal on Scientific Computing* 36.5 (2014), pp. C504–C523.

69. JUNTANG ZHUANG et al. “Adaptive Checkpoint Adjoint Method for Gradient Estimation in Neural ODE.” In: *Proceedings of machine learning research* 119 (2020), pp. 11639–11649.